

2.1

In this question, n is large and p is small; therefore, the distribution should be Poisson.

From the goodness-of-fit summary, we accept the null hypothesis that the underlying process is Poisson. The table comparison and rootgram plot can also confirm this.

```
> #check goodness of Poisson fit
> summary(f)

Goodness-of-fit test for poisson distribution

              X^2 df  P(> X^2)
Likelihood Ratio 1.054648 2 0.5901822

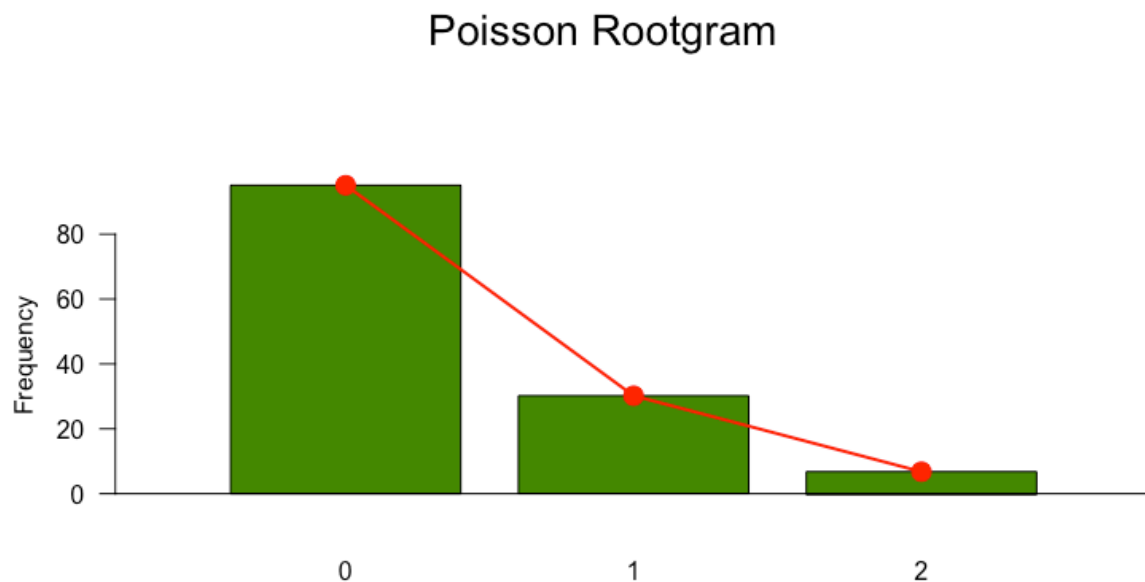
> rootogram(f, main="Poisson Rootgram", xlab = "", ylab="Frequency", rect_gp = gpar(fill =
"chartreuse4"))

> #simulate many Poisson trials using the fitted lambda parameter
> lambda <- f$par

> simulated = rpois(trial_size, lambda[[1]])

> table(l)
l
  0    1    2    3
9053 894  51    2

> table(simulated)
simulated
  0    1    2
9056 893  51
```

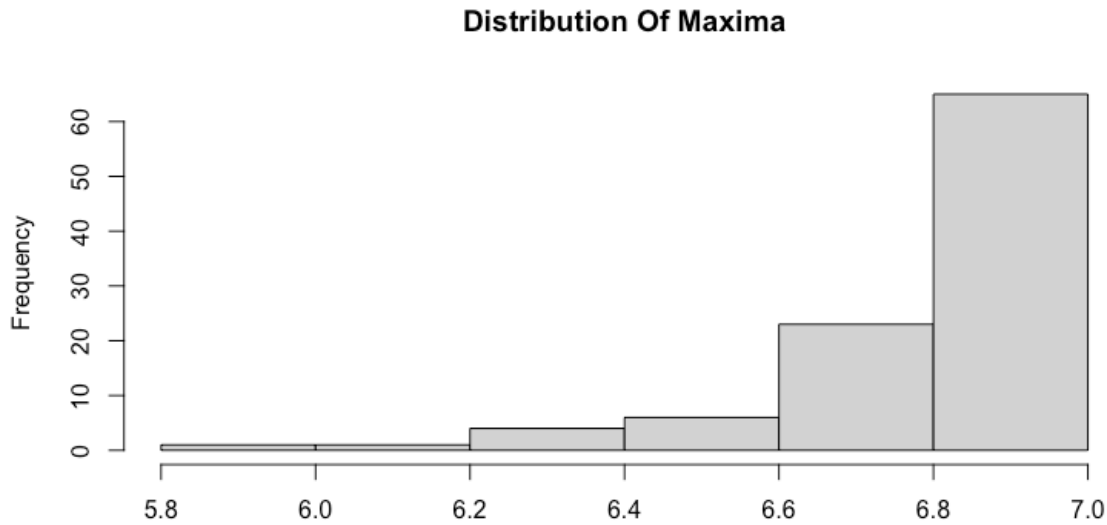


R Code:

```
#2.1
#generate data
sequence_length <- 1000
mutation_rate <- 10^-4
trial_size <- 10000
l <- replicate(trial_size,{sum(rbinom(sequence_length,1,mutation_rate))})
#fit data
library("vcd")
f <- goodfit(l, "poisson")
#check goodness of Poisson fit
summary(f)
rootogram(f, main="Poisson Rootgram",xlab = "",ylab="Frequency", rect_gp = gpar(fill =
"chartreuse4"))

#simulate many Poisson trials using the fitted lambda parameter
lambda <- f$par
simulated = rpois(trial_size,lambda[[1]])
table(l)
table(simulated)
```

2.2



Based on the plot above, it is evident that the maximum likelihood estimation for $\hat{\theta}$ is 7 for the maximum of 25 independent identically distribution uniform random variable.

The theoretical justification is:

Let X_i be an independent identically distributed uniform random variable, $Unif(a, b), a < b$, and $Y_n = \max(X_1, X_2, \dots, X_n)$

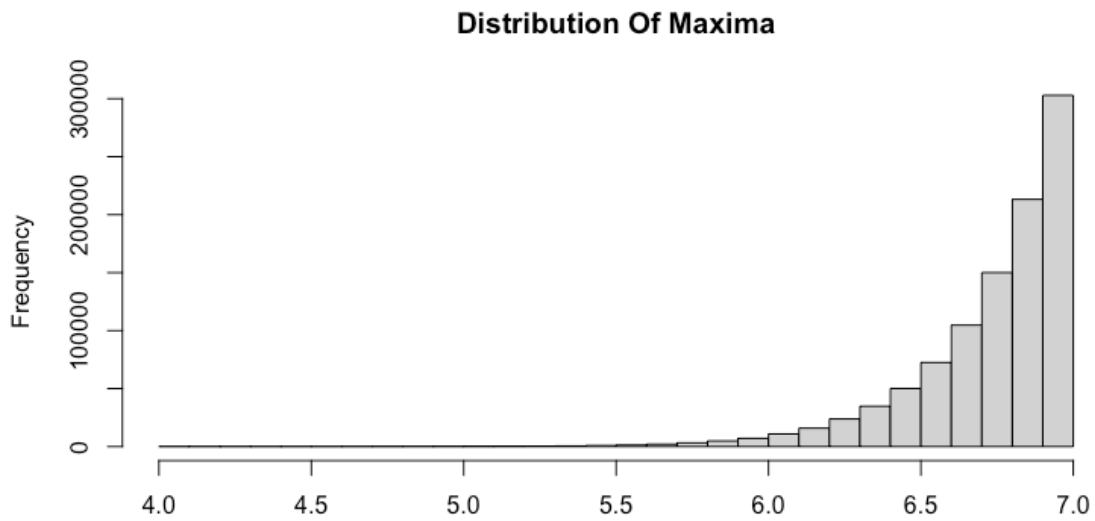
$$p(Y_n \leq x) = p(X_1 \leq x \& X_2 \leq x \& \dots \& X_n \leq x) = \left(\frac{x-a}{b-a}\right)^n$$

Let δ be an infinitesimally small number:

$$\lim_{n \rightarrow \infty} p(Y_n \leq b - \delta) = \left(\frac{b-a-\delta}{b-a}\right)^n = 0$$

Therefore, as n increases, the maximum converges to $b = 7$

Let's increase the number of trials B from 100 to 10000, and plot the distribution of maxima again:



R Code:

#2.2

#generate data

```
generator <- function(n=25,min=0,max=7){
```

```
  return(max(runif(n,0,7)))
```

```
}
```

B = 100

```
l <- replicate(B,generator())
```

#plot data

```
hist(l,xlab="",main="Distribution Of Maxima")
```

2.3

a.

The fact that 20 amino acids have redundant expressions due to 64 codon spellings is verified below.

```
> table(mtb$AmAcid)
```

Ala	Arg	Asn	Asp	Cys	End	Gln	Glu	Gly	His	Ile	Leu	Lys	Met	Phe	Pro	Ser	Thr	Trp	Tyr
4	6	2	2	2	3	2	2	4	2	3	6	2	1	2	4	6	4	1	2

Val
4

```
> table(mtb$Codon)
```

AAA	AAC	AAG	AAT	ACA	ACC	ACG	ACT	AGA	AGC	AGG	AGT	ATA	ATC	ATG	ATT	CAA	CAC	CAG	CAT
1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
CCA	CCC	CCG	CCT	CGA	CGC	CGG	CGT	CTA	CTC	CTG	CTT	GAA	GAC	GAG	GAT	GCA	GCC	GCG	GCT
1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
GGA	GGC	GGG	GGT	GTA	GTC	GTG	GTT	TAA	TAC	TAG	TAT	TCA	TCC	TCG	TCT	TGA	TGC	TGG	TGT
1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
TTA	TTC	TTG	TTT																
1	1	1	1																

b.

The “PerThous” variable refers to the frequency that a codon would appear every thousand codons.

It can be computed from the command
“(mtb\$Number/sum(mtb\$Number))*1000”

c.

the strongest bias belongs to the isoleucine amino acid. There are three codon spellings for isoleucine and the greatest bias for ATC is 46.3%.

```

> #c
> library(dplyr)

> bias_transform <- function(t = mtb){
+   new_mtb <- t %>%
+     group_by(AmAcid) %>%
+     mutate(freq = Number/sum(Number), redundant = length(factor(Codon)) [TRUNCATED])

> bias_transform()
# A tibble: 1 x 7
# Groups:   AmAcid [1]
  AmAcid Codon Number PerThous freq redundant bias
  <chr>   <chr>   <int>   <dbl> <dbl>      <int> <dbl>
1 Ile    ATC      45551    33.9 0.796         3 0.463

```

R Code:

#2.3

#a

```
mtb = read.table("~/Desktop/M_tuberculosis.txt",header=TRUE)
```

```
table(mtb$AmAcid)
```

```
table(mtb$Codon)
```

#b

```
(mtb$Number/sum(mtb$Number))*1000
```

#c

```
library(dplyr)
```

```
bias_transform <- function(t = mtb){
```

```
  new_mtb <- t %>%
```

```
  group_by(AmAcid) %>%
```

```
  mutate(freq = Number/sum(Number), redundant = length(factor(Codon))) %>%
```

```
  mutate(bias = abs(freq-(1/redundant)))
```

```
  return(new_mtb[which.max(new_mtb$bias),])}
```

```
bias_transform()
```

2.4

Question 2.4 was completed with help from a tutorial on “Biostring” posted by the Stanford University:

https://web.stanford.edu/class/bios221/labs/biostrings/lab_1_biostrings.html

a.


To see the complete sequence, use the “as.character(staph[i])” expression.

```
> #a
> staph = readDNAStringSet("~/Desktop/staphsequence.ffn.txt", "fasta")

> staph[1:3]
DNAStringSet object of length 3:
      width seq                                     names
[1]  1362 ATGTCGGAAAAAGAAATTTGGGAAA...GAAAAAGAAATAAGAAATGTATAA lcl|INC_002952.2_c...
[2]  1134 ATGATGGAATTCATTTAAAAGAG...ATTTTACCAATCAGAACTTACTAA lcl|INC_002952.2_c...
[3]   246 GTGATTATTTTGGTTCAAGAAGTTG...ATCATTATCAAGGTGAACAATGA lcl|INC_002952.2_c...
```

b.

Herein we use the built-in function “alphabetFrequency” from the package “Biostring” for fixed window analysis



C.

Herein we use the built-in function

“letterFrequencyInSlidingWindow” from the package “Biostring”
for sliding window analysis

```
58 #b
59 window <- 100
60 GC_content <- letterFrequencyInSlidingView(staph[[1]], window, c("G","C"))
61 GC_content
```

50:18 (Top Level) R Script

Console Terminal x

~ /

```
> #b
> window <- 100
> GC_content <- letterFrequencyInSlidingView(staph[[1]], window, c("G","C"))
> GC_content
      G C
[1,] 18 14
[2,] 18 15
[3,] 19 15
[4,] 18 15
[5,] 18 15
[6,] 18 15
[7,] 17 15
[8,] 16 15
[9,] 16 15
[10,] 17 15
[11,] 17 15
[12,] 17 15
[13,] 18 15
[14,] 18 15
[15,] 18 15
[16,] 19 15
[17,] 19 15
[18,] 19 15
[19,] 20 15
[20,] 20 16
[21,] 19 16
[22,] 18 16
[23,] 17 16
[24,] 17 17
[25,] 18 17
[26,] 18 17
```

```
63 #c
64 GC_fraction <- GC_content/window
65 GC_fraction
```

61:1 (Top Level) R Script

Console Terminal x

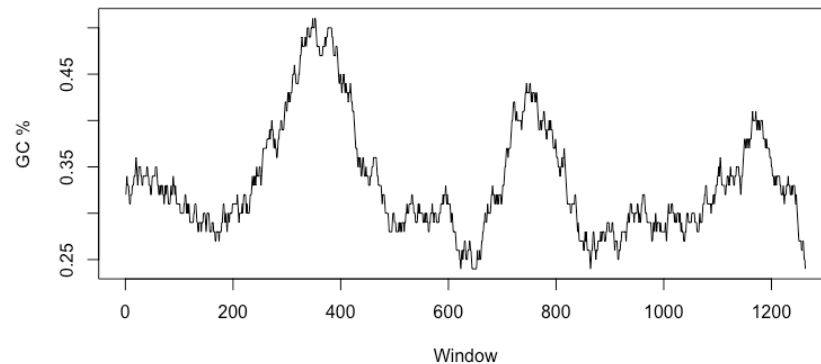
~ /

```
> #c
> GC_fraction <- GC_content/window
> GC_fraction
      G C
[1,] 0.18 0.14
[2,] 0.18 0.15
[3,] 0.19 0.15
[4,] 0.18 0.15
[5,] 0.18 0.15
[6,] 0.18 0.15
[7,] 0.17 0.15
[8,] 0.16 0.15
[9,] 0.16 0.15
[10,] 0.17 0.15
[11,] 0.17 0.15
[12,] 0.17 0.15
[13,] 0.18 0.15
[14,] 0.18 0.15
[15,] 0.18 0.15
[16,] 0.19 0.15
[17,] 0.19 0.15
[18,] 0.19 0.15
[19,] 0.20 0.15
[20,] 0.20 0.16
[21,] 0.19 0.16
[22,] 0.18 0.16
[23,] 0.17 0.16
[24,] 0.17 0.17
[25,] 0.18 0.17
[26,] 0.18 0.17
[27,] 0.18 0.17
[28,] 0.18 0.17
```


d.

We could plot the GC fraction along the window sequence.

Here we can see the plot from part (d) roughly follows that of (b).



R Code:

#2.4

```
library("Biostrings")
```

#a

```
staph = readDNASTringSet("~/Desktop/staphsequence.ffn.txt", "fasta")
```

```
staph[1:3]
```

#b

```
library(Biostrings)
```

```
staph <- readDNASTringSet("~/Desktop/staphsequence.ffn.txt", "fasta")
```

```
window <- 100
```

```
l <- length(staph[[1]])
```

```
start <- (c(1:as.integer(l/window))-1)*window
```

```
end <- start + window
```

```
view <- Views(staph[[1]],start=start,end=end)
```

```
gc_fixed_window <- rowSums(alphabetFrequency(view)[, c(2,3)]/window)
plot(gc_fixed_window, type = 'l')
#c
window <- 100
GC_content <- letterFrequencyInSlidingView(staph[[1]], window, c("G", "C"))
GC_content
GC_fraction <- GC_content/window
GC_fraction


#d
GC_roll <- rowSums(GC_fraction)
plot(GC_roll, type = 'l', ylab="GC %", xlab="Window")
```

2.5

$B(1,1)$ is flat, hence the name “uniform” distribution.

$B(\frac{1}{2},\frac{1}{2})$ is flat at its central region then curves up at its tails.

Using formula $E(X) = \frac{\alpha}{\alpha+\beta}$, we see that $B(\frac{1}{2},\frac{1}{2})$ and $B(1,1)$ have the same mean at 0.5 while the rest of the B s have the same mean at 0.25.



R Code:

```
s <- seq(0,1,by=0.005)
d_10_30 <- dbeta(s, 10, 30)
d_20_60 <- dbeta(s, 20, 60)
d_50_150 <- dbeta(s, 50, 150)
d_1_1 <- dbeta(s, 1, 1)
d_h_h <- dbeta(s, 1/2, 1/2)

plot(s,d_10_30,type="l",lty=2,ylab="",ylim=c(0,15),col="red")
lines(s,d_20_60,col="green",lty=2)
lines(s,d_50_150,col="blue",lty=2)
lines(s,d_1_1, col="purple",lty=2)
```

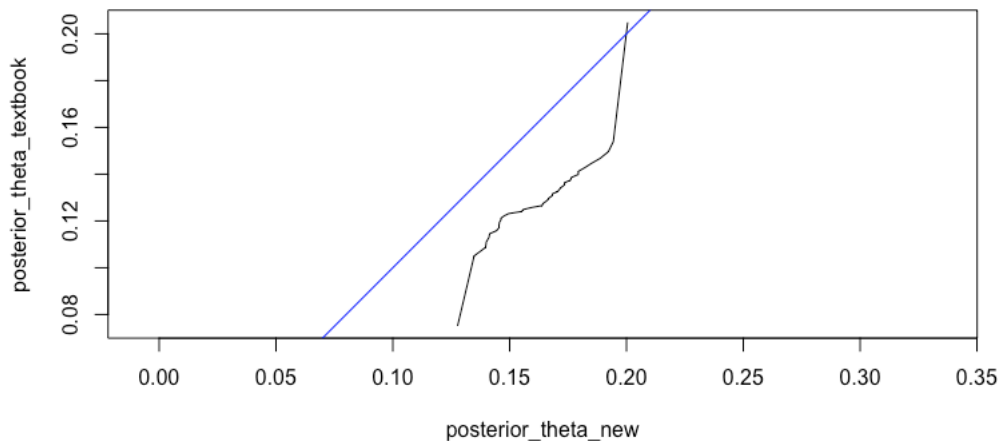
```
lines(s,d_h_h, col="orange",lty=2)
```

```
legend( "topright",  
c("(10,30)","(20,60)","(50,150)","(1,1)","(1/2,1/2)"),text.col=c("red","green","blue","purple","orange"))
```

2.6

The prior distribution for the textbook example is $B(50,350)$ and its posterior distribution given $n = 300, Y = 40$ is $B(90,610)$.

I chose the prior distribution to be $B(20,50)$ and plotted its posterior against that from the textbook example above in a QQ plot.



As expected, they are not matched because the theoretical posterior distribution should be $B(60,310)$, quite different from $B(90,610)$.

R Code:

```
#2.6
#the posterior distribution from textbook where alpha = 50, beta = 350
posterior_theta_textbook = rbeta(n = 1e6, 90, 610)
#my own posterior distribution generated from alpha = 20, beta= 50
rtheta = rbeta(100000, 20, 50)
y = vapply(rtheta, function(th) {rbinom(1, prob = th, size = 300)}, numeric(1))
posterior_theta_new = rtheta[ y == 40 ]
qqplot(posterior_theta_new, posterior_theta_textbook, type = "l", asp = 1)
abline(a = 0, b = 1, col = "blue")
```